

Vector Clock and Singhal-Kshemkalyani Optimization: Project Report

1. Introduction

This report presents the implementation and comparative analysis of two distributed clock synchronization algorithms: the basic Vector Clock algorithm and its optimized variant, the Singhal-Kshemkalyani optimization. The goal is to demonstrate the strong consistency property of vector clocks and measure the performance improvements achieved through optimization.

1.1 Objectives

- Implement basic Vector Clock algorithm
- Implement Singhal-Kshemkalyani optimization
- Demonstrate strong consistency properties
- Compare performance overheads and message communication savings
- Analyze scalability with varying process counts

1.2 Problem Statement

In a distributed system of n nodes connected in a graph topology, each process executes three types of events:

- **Internal events:** Process-local operations
- **Send events:** Message transmission to neighbors
- **Receive events:** Message reception from neighbors

The challenge is to maintain causal ordering while minimizing communication overhead.

2. Algorithm Implementation

2.1 Basic Vector Clock Algorithm

The basic Vector Clock algorithm maintains a vector of n logical clocks, one for each process in the system.

2.1.1 Algorithm Details

```
// Vector Clock Operations
void increment_clock() {
    vector_clock[process_id - 1]++;
}

void update_clock(const vector<int>& received_clock) {
```

```

    for (int i = 0; i < num_processes; i++) {
        vector_clock[i] = max(vector_clock[i], received_clock[i]);
    }
    increment_clock();
}

```

2.1.2 Message Structure

- **Full Vector Clock:** Sends complete vector clock with every message
- **Message Size:** $O(n)$ entries per message
- **Overhead:** Linear growth with system size

2.2 Singhal-Kshemkalyani Optimization

The Singhal-Kshemkalyani optimization reduces message overhead by sending only changed clock entries.

2.2.1 Optimization Strategy

```

vector<pair<int, int>> get_optimized_clock_entries(int target_process)
{
    vector<pair<int, int>> entries;

    for (int i = 0; i < num_processes; i++) {
        int process_id_1_indexed = i + 1;
        if (vector_clock[i] > last_sent_clock[i]) {
            entries.push_back({process_id_1_indexed,
vector_clock[i]});
        }
    }

    last_sent_clock = vector_clock;
    return entries;
}

```


2.2.2 Key Features

- **Selective Transmission:** Only changed entries are sent
 - **Last Sent Tracking:** Maintains history of sent values
 - **Message Reduction:** Significant reduction in message size
-

3. System Design and Implementation

3.1 Architecture Overview

The system implements a multi-threaded simulation where each process runs as a separate thread:



3.2 Event Generation Model

3.2.1 Exponential Distribution

Events are generated with exponentially distributed inter-event times: - **Parameter:** λ (lambda) - average inter-event time - **Distribution:** Exponential with mean $1/\lambda$

3.2.2 Event Type Selection

Event types are selected probabilistically based on parameter α : - **Internal Events:** Probability = $\alpha/(\alpha+1)$ - **Send Events:** Probability = $1/(\alpha+1)$

3.3 Threading and Synchronization

3.3.1 Thread Safety

- **Message Queues:** Mutex-protected shared queues

- **Logging:** Thread-safe file writing
- **Statistics:** Atomic counters for performance metrics

3.3.2 Communication Model

- **Message Passing:** Queue-based inter-process communication
 - **Neighbor Selection:** Random selection from adjacency list
 - **Event Processing:** Non-blocking message reception
-

4. Strong Consistency Demonstration

4.1 Causal Ordering Properties

The implementation demonstrates the strong consistency property of vector clocks:

Property 1: If two events x and y have timestamps vh and vk, respectively, then: - $x \rightarrow y$ (x causally precedes y) $\Leftrightarrow vh < vk$

Property 2: If two events x and y have timestamps vh and vk, respectively, then: - $x \parallel y$ (x and y are concurrent) $\Leftrightarrow vh \parallel vk$

4.2 Event Logging Format

All events are logged with both real-time and vector timestamps:

```
Process1 internal event e11 at 10:00, vc: [1 0 0 0]
Process2 internal event e21 at 10:01, vc: [0 1 0 0]
Process3 send event m31 to process2 at 10:02, vc: [0 0 1 0]
Process2 receive event m31 from process3 at 10:05, vc: [0 3 1 0]
```

4.3 Consistency Verification

The logs demonstrate proper causal ordering: - **Internal events:** Increment local clock component - **Send events:** Include current vector clock in message - **Receive events:** Merge received clock with local clock

5. Performance Analysis

5.1 Experimental Setup

5.1.1 Test Configuration

- **Process Count:** 10-15 processes (increments of 1)
- **Lambda (λ):** 5ms (exponential distribution parameter)
- **Alpha (α):** 1.5 (internal to send event ratio)
- **Messages per Process:** 50
- **Topology:** Ring topology for scalability
- **Test Runs:** Multiple executions for statistical significance

5.1.2 Performance Metrics

- **Message Overhead:** Average entries per message
- **Total Messages:** Number of messages sent
- **Execution Time:** Runtime performance
- **Space Utilization:** Memory usage per process

5.2 Results and Analysis

5.2.1 Message Overhead Comparison

| Processes | VC Avg Entries | SK Avg Entries | Improvement |
|-----------|----------------|----------------|-------------|
| 10 | 10.0 | 1.0 | 90.0% |
| 11 | 11.0 | 1.0 | 90.9% |
| 12 | 12.0 | 1.0 | 91.7% |
| 13 | 13.0 | 1.0 | 92.3% |
| 14 | 14.0 | 1.0 | 92.9% |
| 15 | 15.0 | 1.0 | 93.3% |

Average Improvement: 91.85%

5.2.2 Key Observations

1. **Consistent Optimization:** SK optimization consistently reduces message size by ~90-93%
2. **Scalability:** Improvement increases with process count

3. **Fixed Overhead:** SK maintains constant 1 entry per message regardless of process count
4. **Linear Growth:** VC overhead grows linearly with process count

5.2.3 Execution Time Analysis

| Processes | VC Time (s) | SK Time (s) | Time Ratio |
|-----------|-------------|-------------|------------|
| 10 | 1.15 | 1.26 | 1.10 |
| 11 | 1.14 | 1.05 | 0.92 |
| 12 | 1.16 | 1.13 | 0.97 |
| 13 | 1.16 | 1.07 | 0.92 |
| 14 | 1.14 | 1.14 | 1.00 |
| 15 | 1.12 | 1.25 | 1.12 |

Average Time Ratio: 1.01 (SK is approximately 1% slower on average)


5.3 Scalability Analysis

5.3.1 Message Overhead Trends

The analysis reveals clear scalability patterns:

1. **Vector Clock:** Message overhead grows linearly with process count
 - Formula: $O(n)$ entries per message
 - Impact: Network bandwidth grows quadratically with system size
2. **Singhal-Kshemkalyani:** Message overhead remains constant
 - Formula: $O(1)$ entries per message (average case)
 - Impact: Network bandwidth grows linearly with system size

The below graphs represent different metrics confirming on easy scalability for **Singhal-Kshemkalyani algorithm**.



5.4 Anomaly Analysis

5.4.1 Expected Patterns

- Consistent Improvement:** SK consistently outperforms VC
- Scalable Benefits:** Improvement increases with system size
- Fixed Overhead:** SK maintains constant message size

5.4.2 Observed Anomalies

- Execution Time Variance:** Some variation in execution times due to:
 - Random event generation
 - System load variations
 - Thread scheduling differences

5.4.3 Statistical Significance

- Consistent Results:** Multiple test runs show consistent patterns
- Clear Trends:** Improvement percentage shows clear upward trend
- Reliable Metrics:** Message overhead measurements are deterministic

6. Algorithm Complexity Analysis

6.1 Time Complexity

6.1.1 Vector Clock

- **Internal Event:** $O(1)$ - single increment operation
- **Send Event:** $O(n)$ - copy entire vector clock
- **Receive Event:** $O(n)$ - merge n clock components
- **Overall:** $O(n)$ per message

6.1.2 Singhal-Kshemkalyani

- **Internal Event:** $O(1)$ - single increment operation
- **Send Event:** $O(n)$ - check all components for changes
- **Receive Event:** $O(k)$ - merge k changed components ($k \leq n$)
- **Overall:** $O(n)$ worst case, $O(1)$ average case

6.2 Space Complexity

6.2.1 Vector Clock

- **Per Process:** $O(n)$ - vector clock storage
- **Per Message:** $O(n)$ - full vector clock transmission
- **Total:** $O(n^2)$ for n processes

6.2.2 Singhal-Kshemkalyani

- **Per Process:** $O(n)$ - vector clock + last sent tracking
- **Per Message:** $O(k)$ - only changed components ($k \leq n$)
- **Total:** $O(n^2)$ worst case, $O(n)$ average case

6.3 Communication Complexity

6.3.1 Vector Clock

- **Message Size:** $O(n)$ entries per message
- **Total Overhead:** $O(n^2)$ for n messages
- **Bandwidth:** Grows quadratically with system size

6.3.2 Singhal-Kshemkalyani

- **Message Size:** $O(1)$ entries per message (average)
- **Total Overhead:** $O(n)$ for n messages

- **Bandwidth:** Grows linearly with system size
-

7. Implementation Challenges and Solutions

7.1 Technical Challenges

7.1.1 Thread Synchronization

Challenge: Ensuring thread-safe communication between processes **Solution:** - Mutex-protected message queues - Atomic counters for statistics - Thread-safe logging with file locks

7.1.2 Event Generation

Challenge: Generating realistic event patterns **Solution:** - Exponential distribution for timing - Probabilistic event type selection - Random neighbor selection

7.1.3 Performance Measurement

Challenge: Accurate measurement of message overhead **Solution:** - Atomic counters for statistics - Precise timing measurements - Comprehensive logging system

7.2 Algorithmic Challenges

7.2.1 Singhal-Kshemkalyani Implementation

Challenge: Correctly tracking last sent values **Solution:** - Maintain separate last_sent_clock vector - Update tracking after each send operation - Ensure proper initialization

7.2.2 Termination Condition

Challenge: Ensuring all processes complete properly **Solution:** - Global termination flag - Message count tracking - Proper thread synchronization

8. Conclusions and Future Work

8.1 Key Findings

1. **Significant Performance Improvement:** Singhal-Kshemkalyani optimization achieves 91.85% average reduction in message overhead
2. **Scalable Benefits:** Performance improvement increases with system size

3. **Correctness Preservation:** Optimization maintains all vector clock consistency properties
4. **Practical Impact:** Substantial bandwidth and network efficiency gains

8.2 Theoretical Validation

The experimental results validate theoretical expectations: - **Linear Growth:** VC overhead grows linearly with process count - **Constant Overhead:** SK maintains constant message size - **Causal Ordering:** Both algorithms preserve causal relationships

8.3 Practical Implications

1. **Network Efficiency:** 90%+ reduction in network traffic
2. **Scalability:** Better performance in large distributed systems
3. **Resource Utilization:** Reduced bandwidth and processing overhead
4. **Cost Benefits:** Lower infrastructure costs for large deployments

8.4 Future Work

8.4.1 Algorithmic Improvements

- **Adaptive Optimization:** Dynamic adjustment based on system load
- **Hybrid Approaches:** Combine multiple optimization techniques
- **Machine Learning:** Predict optimal message patterns

8.4.2 System Enhancements

- **Real Network:** Test on actual distributed systems
- **Fault Tolerance:** Handle process failures and network partitions
- **Dynamic Topology:** Support for changing network topologies

8.4.3 Performance Analysis

- **Larger Scale:** Test with hundreds/thousands of processes
- **Different Topologies:** Analyze performance across various network structures
- **Real-world Workloads:** Test with actual application patterns

9. Appendices

Appendix A: Source Code Structure

```
vector_clocks/
├── VC-CS25RESCH04001.cpp          # Basic Vector Clock implementation
└── SK-CS25RESCH04001.cpp          # Singhal-Kshemkalyani optimization
```

```
└── performance_comparison.cpp # Performance testing script
└── generate_graphs.py       # Graph generation
└── Makefile                  # Build automation
└── readme.txt                # Execution instructions
└── inp-params.txt           # Input parameters
```

Appendix B: Compilation Instructions

```
# Compile all programs
make all

# Run Vector Clock
./VC-CS25RESCH04001

# Run Singhal-Kshemkalyani
./SK-CS25RESCH04001

# Run performance comparison
./performance_comparison

# Generate graphs
python3 generate_graphs.py
```

Appendix C: Input Format

```
n λ α m
1 neighbor1 neighbor2 ...
2 neighbor1 neighbor2 ...
...
n neighbor1 neighbor2 ...
```

Appendix D: Sample Output

```
Process1 internal event e11 at 10:00, vc: [1 0 0 0]
Process2 internal event e21 at 10:01, vc: [0 1 0 0]
Process3 send event m31 to process2 at 10:02, vc: [0 0 1 0]
Process2 receive event m31 from process3 at 10:05, vc: [0 3 1 0]
```
